

OpenStack Clients

Guide

trunk (May 13, 2013)



OpenStack Clients Guide

trunk (2013-05-13)

Copyright © 2009-2013 OpenStack Foundation All rights reserved.

The OpenStack clients are command-line interfaces (CLIs) that let you run simple commands to make OpenStack API calls. These open-source Python clients are easy to learn and use. The OpenStack APIs are RESTful APIs that use all aspects of the HTTP protocol, including methods, URIs, media types, and response codes. To request OpenStack services, you must first issue an authentication request to the OpenStack Identity Service v2.0.

Licensed under the Apache License, Version 2.0 (the "License"); you may not use this file except in compliance with the License. You may obtain a copy of the License at

<http://www.apache.org/licenses/LICENSE-2.0>

Unless required by applicable law or agreed to in writing, software distributed under the License is distributed on an "AS IS" BASIS, WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied. See the License for the specific language governing permissions and limitations under the License.

List of Tables

1.1. OpenStack Clients	1
2.1. OpenStack Clients Prerequisite Software	3
3.1. OpenStack Clients Environment Variables	6
9.1. List of configuration flags for NFS	31

List of Examples

9.1. nova Usage	46
9.2. Positional arguments	46
9.3. Optional arguments	50

1. OpenStack Clients Overview

The OpenStack clients are command-line interfaces (CLIs) that let you run simple commands to make OpenStack API calls. These open-source Python clients are easy to learn and use.

The OpenStack APIs are RESTful APIs that use all aspects of the HTTP protocol, including methods, URIs, media types, and response codes. Internally, each client command runs cURL commands that embed API requests.

To install an OpenStack client, see [Chapter 2, “Install the OpenStack Clients”](#) [3].

To request OpenStack services either through the clients or through the APIs directly, you must first issue an authentication request to the OpenStack Identity Service v2.0. To do so, you can run the **credentials** command, which is a nova client command.

For example, to use the OpenStack Compute API from the command line, complete these steps:

1. Install the nova client.
2. Issue the nova **credentials** command.
3. Issue other nova client commands, such as **boot**, **list**, and so on.



Note

An OpenStack common client is in development.

To install the OpenStack clients on a Mac OS X or Linux system, use **pip** because it is easy and ensures that you get the latest version of the client from the [Python Package Index](#). Also, **pip** lets you update or remove a package later on.

Use the following OpenStack clients to access the CLIs:

Table 1.1. OpenStack Clients

Client	Associated API	Description	See
cinder	Block storage	Create and delete volumes, attach volumes to and detach volumes from servers, create and delete snapshots, create volumes from snapshots, and get volume statistics.	Chapter 10, “cinder Client” [53]
glance	Image service	Manage images. For example, add and set permissions on images.	Chapter 7, “glance Client” [12]
keystone	Identity service	Create and manage users, tenants, roles, endpoints, and credentials.	Chapter 6, “keystone Client” [10]
nova	Compute, Compute extensions	Authenticate, launch servers, set security groups, control IP addresses on servers, control volumes, and create images.	Chapter 9, “nova Client” [28]
quantum	Networking	Configure networks for guest servers.	Chapter 8, “quantum Client” [18]
swift	Object storage	Gather statistics, list items, update metadata, upload, download and delete files stored by the object storage service. Provides access to a swift installation for ad hoc processing.	Chapter 11, “swift Client” [55]

Client	Associated API	Description	See
heat	Orchestration	Launch stacks from templates, view details of running stacks (including events and resources), update and delete stacks.	Chapter 12, "heat Client" [57]

2. Install the OpenStack Clients

Table of Contents

Before You Begin	3
Install the OpenStack Clients	4

To manage your servers, images, volumes, isolated networks, and other cloud resources from the command line, install and use the open-source clients.

To install the clients, first install some prerequisite software. Then, install the Python packages. Each Python package is an OpenStack client.

Before You Begin

Before you begin, install the following prerequisite software:

Table 2.1. OpenStack Clients Prerequisite Software

Prerequisite	Description
Python 2.6 or later	Currently, the clients do not support Python 3.
setuptools package	<p>Installed by default on Mac OS X.</p> <p>Many Linux distributions provide packages to make setuptools easy to install.</p> <p>Search your package manager for setuptools to find an installation package. If you cannot find one, download the setuptools package directly from http://pypi.python.org/pypi/setuptools.</p>
pip package	<p>To install the clients on a Mac OS X or Linux system, use pip. It is easy to use and ensures that you get the latest version of the clients from the Python Package Index. Also, it lets you update or remove the packages later on.</p> <p>Install pip through the package manager for your system:</p> <ul style="list-style-type: none">• Mac OS X<pre>\$ sudo easy_install pip</pre>• Ubuntu 12.04<p>A packaged version enables you to use dpkg or aptitude to install the python-novaclient.</p><pre>\$ aptitude install python-novaclient</pre>• Ubuntu<pre>\$ aptitude install python-pip</pre>• RHEL, CentOS, or Fedora<pre>\$ yum install python-pip</pre>• openSUSE 12.2 and earlier<p>A packaged version available in the Open Build Service enables you to use rpm or zypper to install the the python-novaclient.</p>

Prerequisite	Description
	<pre>\$ zypper install python-pip</pre> <p>Alternatively, you can still use pip.</p> <ul style="list-style-type: none">• openSUSE 12.3 <p>A packaged version enables you to use rpm or zypper to install the python-novaclient.</p> <p>For example:</p> <pre>\$ zypper install python-novaclient</pre>

Install the OpenStack Clients

Procedure 2.1. To install the OpenStack clients:

1. Install or update the client packages

You must install each client separately.

Run the following command to install or update a client package:

```
$ sudo pip install [--update] python-<project>client
```

Where <project> is the project name and is one of the following values:

- nova. Compute API.
- quantum. Networking API.
- keystone. Identity service API.
- glance. Image service API.
- swift. Object storage API.
- cinder. Block storage API.
- heat. Orchestration API.

For example, to install the nova client, run the following command:

```
$ sudo pip install python-novaclient
```

To update the nova client, run the following command:

```
$ sudo pip install --upgrade python-novaclient
```

To remove the nova client, run the following command:

```
$ sudo pip uninstall python-novaclient
```

2. Set environment variables and authenticate

Before you can issue client commands, you must set environment variables and authenticate. See [Chapter 3, “Authenticate” \[5\]](#).

3. Authenticate

To authenticate a tenant to run commands for the OpenStack clients, you issue an authentication request to the OpenStack Identity Service v2.0. To issue this request, you must first set and source environment variables and install the nova client. Then, you issue an authentication request through the nova **credentials** command.

Procedure 3.1. To authenticate:

1. Set environment variables

Before you can issue client commands, you must set environment variables to authenticate.

You can either edit your bash profile to add and set environment variables or use an `openrc` file downloaded from an OpenStack Dashboard.

Either edit your `.bash_profile` file:

```
$ nano ~/.bash_profile
```

Add the following lines to the bash profile. Edit the values for the **OS_USERNAME**, **OS_PASSWORD**, and **OS_TENANT_NAME** variables:

```
export OS_USERNAME=username
export OS_PASSWORD=password
export OS_TENANT_NAME=tenant
export OS_AUTH_URL=https://identity.api.rackspacecloud.com/v2.0/ #an example, insert
your endpoint here
export NOVACLIENT_DEBUG=1
export NOVA_VERSION=2
```

Or download an `openrc` file from the OpenStack Dashboard:

```
#!/bin/bash

# With the addition of Keystone, to use an openstack cloud you should
# authenticate against keystone, which returns a **Token** and **Service
# Catalog**. The catalog contains the endpoint for all services the
# user/tenant has access to - including nova, glance, keystone, swift.
#
# *NOTE*: Using the 2.0 *auth api* does not mean that compute api is 2.0. We
# will use the 1.1 *compute api*
export OS_AUTH_URL=http://10.0.100.102:5000/v2.0

# With the addition of Keystone we have standardized on the term **tenant**
# as the entity that owns the resources.
export OS_TENANT_ID=feacce5a1fc347f88cfc0dee838429d6
export OS_TENANT_NAME=tenant

# In addition to the owning entity (tenant), openstack stores the entity
# performing the action as the **user**.
export OS_USERNAME=username

# With Keystone you pass the keystone password.
echo "Please enter your OpenStack Password: "
read -s OS_PASSWORD_INPUT
export OS_PASSWORD=$OS_PASSWORD_INPUT
```

Source the file:

```
$ source openrc.sh
```

Enter your OpenStack password when prompted.

The following table describes the environment variables:

Table 3.1. OpenStack Clients Environment Variables

Environment Variable	Description
OS_USERNAME	Your OpenStack username.
OS_PASSWORD	Your OpenStack user password.
OS_TENANT_ID	Your tenant ID, usually provided with your username.
OS_TENANT_NAME	Your tenant name, usually provided with your username.
OS_AUTH_URL	The endpoint for the Identity Service (keystone), which the nova client uses for authentication. Include the trailing forward slash (/) in the URL. Otherwise, you receive a 404 error.
NOVACLIENT_DEBUG	Set to 1 to show the underlying cURL commands with embedded API requests in the command responses. Otherwise, omit this variable.
NOVA_VERSION	The version of the API. Set to 2.

After you set the variables, save the file.

2. Set permissions on and source the bash profile

Because the bash profile contains a password, set permissions on it so other people cannot read it:

```
$ chmod 600 ~/.bash_profile
```

To source the variables to make them available in your current shell, run the following command:

```
$ source ~/.bash_profile
```

3. Authenticate

To authenticate, you must install the nova client. To install the nova client, see [Chapter 2, “Install the OpenStack Clients” \[3\]](#).

To verify that you can talk to the API server, authenticate and list images:

```
$ nova credentials
$ nova image-list
```

Then, list networks:

```
$ nova network-list
```

4. Troubleshooting

- If you cannot run commands successfully, make sure that you entered your user name, password, and account number correctly in the bash profile file.
- If you change any environment variables, either log out and back in or source your bash profile again.

- To override some environment variable settings, you can use the options that are listed at the end of the **nova help** output. For example, you can override the `OS_PASSWORD` setting in the bash profile by specifying a password on a nova command, as follows:

```
$ nova --password <password> image-list
```

Where *password* is your password.

4. Get the Version for a Client

Search for the version number.

```
$pip freeze | grep python-  
  
python-glanceclient==0.4.0  
python-keystoneclient==0.1.2  
-e git+https://github.com/openstack/python-novaclient.  
git@077cc0bf22e378c4c4b970f2331a695e440a939f#egg=python_novaclient-dev  
python-quantumclient==0.1.1  
python-swiftclient==1.1.1
```

You can also use the `yolk -l` command to see what version of the CLI you have installed.

```
$yolk -l | grep python-novaclient  
  
python-novaclient - 2.6.10.27      - active development (/Users/your.name/src/  
cloud-servers/src/src/python-novaclient)  
python-novaclient - 2012.1         - non-active
```

5. Get Help for Client Commands

Use the **help** command to get help for commands, parameters, and subcommands for any OpenStack client.

The syntax is:

```
$ <client-name> help
```

For example, to get help for glance client commands, run the following command:

```
$ glance help
```

The **help** command lists the available commands for the specified client.



Note

Depending on your credentials, you might not have permission to use every command.

To get help for a specific command, enter the command name after the **help** command, as follows:

```
$ <client-name> help <command-name>
```

For example, to get help for the glance **image-show** command, enter the following command:

```
$ glance help image-show
```

The **help** command shows the command usage, a description of the command, and descriptions of any positional and optional arguments, as follows:

```
usage: glance image-show [--human-readable] <IMAGE>
```

```
Describe a specific image.
```

```
Positional arguments:
```

```
<IMAGE>          Name or ID of image to describe.
```

```
Optional arguments:
```

```
--human-readable  Print image size in a human-friendly format.
```

6. keystone Client

Table of Contents

keystone Command Reference	10
----------------------------------	----

This chapter describes how to use the keystone client.

To install the client, see [Chapter 2, “Install the OpenStack Clients”](#) [3].

keystone Command Reference

```
usage: keystone [--os-username <auth-user-name>]
               [--os-password <auth-password>]
               [--os-tenant-name <auth-tenant-name>]
               [--os-tenant-id <tenant-id>] [--os-auth-url <auth-url>]
               [--os-region-name <region-name>]
               [--os-identity-api-version <identity-api-version>]
               [--token <service-token>] [--endpoint <service-endpoint>]
               [--os-cacert <ca-certificate>] [--os-cert <certificate>]
               [--os-key <key>] [--insecure] [--username <auth-user-name>]
               [--password <auth-password>] [--tenant_name <tenant-name>]
               [--auth_url <auth-url>] [--region_name <region-name>]
```

catalog	List service catalog, possibly filtered by service.
ec2-credentials-create	Create EC2-compatible credentials for user per tenant
ec2-credentials-delete	Delete EC2-compatible credentials
ec2-credentials-get	Display EC2-compatible credentials
ec2-credentials-list	List EC2-compatible credentials for a user
endpoint-create	Create a new endpoint associated with a service
endpoint-delete	Delete a service endpoint
endpoint-get	Find endpoint filtered by a specific attribute or service type
endpoint-list	List configured service endpoints
role-create	Create new role
role-delete	Delete role
role-get	Display role details
role-list	List all roles
service-create	Add service to Service Catalog
service-delete	Delete service from Service Catalog
service-get	Display service from Service Catalog
service-list	List all services in Service Catalog
tenant-create	Create new tenant
tenant-delete	Delete tenant
tenant-get	Display tenant details
tenant-list	List all tenants
tenant-update	Update tenant name, description, enabled status
token-get	Display the current user token
user-create	Create new user

user-delete	Delete user
user-get	Display user details.
user-list	List users
user-password-update	Update user password
user-role-add	Add role to user
user-role-list	List roles granted to a user
user-role-remove	Remove role from user
user-update	Update user's name, email, and enabled status
discover	Discover Keystone servers and show authentication protocols and
bash-completion	Prints all of the commands and options to stdout.
help	Display help about this program or one of its
subcommands.	

```

Optional arguments:
--os-username <auth-user-name>
    Defaults to env[OS_USERNAME]
--os-password <auth-password>
    Defaults to env[OS_PASSWORD]
--os-tenant-name <auth-tenant-name>
    Defaults to env[OS_TENANT_NAME]
--os-tenant-id <tenant-id>
    Defaults to env[OS_TENANT_ID]
--os-auth-url <auth-url>
    Defaults to env[OS_AUTH_URL]
--os-region-name <region-name>
    Defaults to env[OS_REGION_NAME]
--os-identity-api-version <identity-api-version>
    Defaults to env[OS_IDENTITY_API_VERSION] or 2.0
--token <service-token>
    Defaults to env[SERVICE_TOKEN]
--endpoint <service-endpoint>
    Defaults to env[SERVICE_ENDPOINT]
--os-cacert <ca-certificate>
    Defaults to env[OS_CA_CERT]
--os-cert <certificate>
    Defaults to env[OS_CERT]
--os-key <key>
    Defaults to env[OS_KEY]
--insecure
    Explicitly allow keystoneclient to perform "insecure"
    SSL (https) requests. The server's certificate will
    not be verified against any certificate authorities.
    This option should be used with caution.
--username <auth-user-name>
    Deprecated
--password <auth-password>
    Deprecated
--tenant_name <tenant-name>
    Deprecated
--auth_url <auth-url>
    Deprecated
--region_name <region-name>
    Deprecated

```

7. glance Client

Table of Contents

List Images	12
Add a New Image	12
Before You Add an Image	13
Upload an image to glance	13
Update an image	13
Managing Images	14
glance Command Reference	14

This chapter describes how to use the glance client.

To install the client, see [Chapter 2, “Install the OpenStack Clients” \[3\]](#).

List Images

To see what images are available to you, use this command:

```
$ glance image-list
```

ID	Name	Status	Server
53b205cc-7abc-46eb-aa60-eabc449b4217	natty-image	ACTIVE	
588d93af-645d-4312-a5b0-81347715a91b	tty-image	ACTIVE	
ac6f83b7-078c-47bd-b4c2-4053282da49e	oneiric-image	ACTIVE	
e110fb7d-2a9e-4da5-923f-5565867ce87a	maverick-image	ACTIVE	

You can also narrow down the list by using `grep` to find only the CentOS images with a command like this:

```
$ glance image-list | grep 'natty'
```

12	natty-server-cloudimg-amd64-kernel	ACTIVE	
13	natty-server-cloudimg-amd64	ACTIVE	

Add a New Image

Adding a new image to your OpenStack cloud.

This process uses the following commands:

- `glance image-create`
- `glance member-create`
- `glance member-list`
- `glance image-show`

Before You Add an Image

Ensure you have created an image that is OpenStack compatible. For details, see the [Image Management chapter](#) in the *OpenStack Compute Administration Manual*.

Upload an image to glance

Assuming you had a CentOS 6.3 image in qcow2 format called centos63.qcow2, the following example will upload it to glance and configure it for public access.

```
$ glance image-create --name centos63-image --disk-format=qcow2 --container-format=raw --is-public=True < ./centos63.qcow2
```

Update an image

To update an image, use the following command:

```
$ glance image-update <image>
```

Where `image` is the name or ID of the image that you want to update. You can use the following optional arguments to modify the following image properties:

Optional arguments:

- `-name <NAME>` Name of image.
- `-disk-format <DISK_FORMAT>`
Disk format of image. Acceptable formats: ami, ari, aki, vhd, vmdk, raw, qcow2, vdi, and iso.
- `-container-format <CONTAINER_FORMAT>`
Container format of image. Acceptable formats: ami, ari, aki, bare, and ovf.
- `-owner <TENANT_ID>` Tenant who should own image.
- `-size <SIZE>` Size of image data (in bytes).
- `-min-disk <DISK_GB>` Minimum size of disk needed to boot image (in gigabytes).
- `-min-ram <DISK_RAM>` Minimum amount of ram needed to boot image (in megabytes).
- `-location <IMAGE_URL>`
URL where the data for this image already resides. For example, if the image data is stored in swift, you could specify 'swift://account:key@example.com/container/obj'.
- `-file <FILE>` Local file that contains disk image to be uploaded during update. Alternatively, images can be passed to the client via stdin.
- `-checksum <CHECKSUM>`
Hash of image data used Glance can use for verification.
- `-copy-from <IMAGE_URL>`
Similar to '-location' in usage, but this indicates that the Glance server should immediately copy the data and store it in its configured image store.
- `-is-public [True|False]`

```

    Make image accessible to the public.
  -is-protected [True|False]
    Prevent image from being deleted.
  -property <key=value>
    Arbitrary property to associate with image. May be
    used multiple times.
  -purge-props    If this flag is present, delete all image properties
                  not explicitly set in the update request. Otherwise,
                  those properties not referenced are preserved.
  -human-readable  Print image size in a human-friendly format.

```

To annotate an image with a property that describes the required VIF model, use the `--property` argument.

For example:

```
$ glance image-update --property hw_vif_model=e1000 f16-x86_64-openstack-sda
```

Valid model values vary depending on the `libvirt_type` setting:

libvirt_type setting	Supported model values
qemu or kvm	<ul style="list-style-type: none"> • virtio • ne2k_pci • pcnet • rtl8139 • e1000
xen	<ul style="list-style-type: none"> • netfront • ne2k_pci • pcnet • rtl8139 • e1000

Requesting an unsupported VIF model causes the guest instance to fail to launch.

Managing Images

Adding images and setting the access to them can be managed in glance, but you can create images by taking a snapshot of a running instance and view available images, set or delete image metadata, and delete an image, using the nova client.

glance Command Reference

```

usage: glance [-version] [-d] [-v] [-k] [--cert-file CERT_FILE]
              [--key-file KEY_FILE] [--os-cacert <ca-certificate-file>]
              [--ca-file OS_CACERT] [--timeout TIMEOUT] [--no-ssl-compression]
              [-f] [--dry-run] [--ssl] [-H ADDRESS] [-p PORT]
              [--os-username OS_USERNAME] [-I OS_USERNAME]
              [--os-password OS_PASSWORD] [-K OS_PASSWORD]
              [--os-tenant-id OS_TENANT_ID] [--os-tenant-name OS_TENANT_NAME]

```

```
[--T OS_TENANT_NAME] [--os-auth-url OS_AUTH_URL] [--N OS_AUTH_URL]
[--os-region-name OS_REGION_NAME] [--R OS_REGION_NAME]
[--os-auth-token OS_AUTH_TOKEN] [--A OS_AUTH_TOKEN]
[--os-image-url OS_IMAGE_URL] [--U OS_IMAGE_URL]
[--os-image-api-version OS_IMAGE_API_VERSION]
[--os-service-type OS_SERVICE_TYPE]
[--os-endpoint-type OS_ENDPOINT_TYPE] [--S OS_AUTH_STRATEGY]
<subcommand> ...
```

Command-line interface to the OpenStack Images API.

Positional arguments:

<subcommand>

add	DEPRECATED! Use image-create instead.
clear	DEPRECATED!
delete	DEPRECATED! Use image-delete instead.
details	DEPRECATED! Use image-list instead.
image-create	Create a new image.
image-delete	Delete specified image(s).
image-download	Download a specific image.
image-list	List images you can access.
image-members	DEPRECATED! Use member-list instead.
image-show	Describe a specific image.
image-update	Update a specific image.
index	DEPRECATED! Use image-list instead.
member-add	DEPRECATED! Use member-create instead.
member-create	Share a specific image with a tenant.
member-delete	Remove a shared image from a tenant.
member-images	DEPRECATED! Use member-list instead.
member-list	Describe sharing permissions by image or tenant.
members-replace	DEPRECATED!
show	DEPRECATED! Use image-show instead.
update	DEPRECATED! Use image-update instead.
help	Display help about this program or one of its subcommands.

Optional arguments:

--version	show program's version number and exit
-d, --debug	Defaults to env[GLANCECLIENT_DEBUG]
-v, --verbose	Print more verbose output
-k, --insecure	Explicitly allow glanceclient to perform "insecure SSL" (https) requests. The server certificate will not be verified against any certificate authorities. This option should be used with caution.
--cert-file CERT_FILE	Path of certificate file to use in SSL connection. This file can optionally be prepended with the private key.
--key-file KEY_FILE	Path of client key to use in SSL connection. This option is not necessary if your key is prepended to your cert file.

```
-os-cacert <ca-certificate-file>
    Path of CA TLS certificate(s) used to verify the remote
    server's certificate. Without this option glance looks
    for the default system CA certificates.
-ca-file OS_CACERT DEPRECATED! Use -os-cacert.
-timeout TIMEOUT    Number of seconds to wait for a response
-no-ssl-compression Disable SSL compression when using https.
-f, -force          Prevent select actions from requesting user
                    confirmation.
-dry-run            DEPRECATED! Only used for deprecated legacy commands.
-ssl                DEPRECATED! Send a fully-formed endpoint using -os-
                    image-url instead.
-H ADDRESS, -host ADDRESS
                    DEPRECATED! Send a fully-formed endpoint using -os-
                    image-url instead.
-p PORT, -port PORT DEPRECATED! Send a fully-formed endpoint using -os-
                    image-url instead.
-os-username OS_USERNAME
                    Defaults to env[OS_USERNAME]
-I OS_USERNAME      DEPRECATED! Use -os-username.
-os-password OS_PASSWORD
                    Defaults to env[OS_PASSWORD]
-K OS_PASSWORD      DEPRECATED! Use -os-password.
-os-tenant-id OS_TENANT_ID
                    Defaults to env[OS_TENANT_ID]
-os-tenant-name OS_TENANT_NAME
                    Defaults to env[OS_TENANT_NAME]
-T OS_TENANT_NAME   DEPRECATED! Use -os-tenant-name.
-os-auth-url OS_AUTH_URL
                    Defaults to env[OS_AUTH_URL]
-N OS_AUTH_URL      DEPRECATED! Use -os-auth-url.
-os-region-name OS_REGION_NAME
                    Defaults to env[OS_REGION_NAME]
-R OS_REGION_NAME   DEPRECATED! Use -os-region-name.
-os-auth-token OS_AUTH_TOKEN
                    Defaults to env[OS_AUTH_TOKEN]
-A OS_AUTH_TOKEN, -auth_token OS_AUTH_TOKEN
                    DEPRECATED! Use -os-auth-token.
-os-image-url OS_IMAGE_URL
                    Defaults to env[OS_IMAGE_URL]
-U OS_IMAGE_URL, -url OS_IMAGE_URL
                    DEPRECATED! Use -os-image-url.
-os-image-api-version OS_IMAGE_API_VERSION
                    Defaults to env[OS_IMAGE_API_VERSION] or 1
-os-service-type OS_SERVICE_TYPE
                    Defaults to env[OS_SERVICE_TYPE]
-os-endpoint-type OS_ENDPOINT_TYPE
                    Defaults to env[OS_ENDPOINT_TYPE]
-S OS_AUTH_STRATEGY, -os_auth_strategy OS_AUTH_STRATEGY
                    DEPRECATED! This option is completely ignored.
```

See "glance help COMMAND" for help on a specific command.

8. quantum Client

Table of Contents

Overview	18
Argument parts of API 2.0 command	18
Features from cliff	18
Features from API	21
Sample quantum command	22
quantum Client Reference	25

This chapter describes how to use the quantum client.

To install the client, see [Chapter 2, “Install the OpenStack Clients”](#) [3].

Overview

Argument parts of API 2.0 command

In general, quantum client command arguments divide into three parts:

Known options

These options are listed in the command's help usage text.

Positional arguments

Positional arguments are mandatory information for an API resource. They must be given in the order.

Unknown options

Unknown options are complementary to known options. To define an unknown option, the format is `--optionname [type=int|bool|dict...] [list=true] [optionvalue]*`. There can be multiple option values for a certain optionname. When there is no optionvalue given, the option is regarded as a `bool` one and value is `true`. The type is python built-in type, such as `int`, `bool`, `float` and `dict`, defaulted to `string` if not given. Unknown options can be used to provide values for creating, updating a resource and to provide filters to list resources. It is also useful to implement API extension when the known options are not included in the command. pseudo argument `--` can be used if the options after it need take advantage of unknown options parser.

Features from cliff

Interactive mode

If there is no command specified, the quantum client will enter into interactive mode:

```

$quantum --os-username admin --os-password password --os-tenant-name admin --
os-auth-url http://localhost:5000/v2.0
(quantum) help

Shell commands (type help <topic>):
=====
cmdenvironment  edit  hi      l    list  pause  r      save  shell      show
ed              help  history li   load  py     run   set    shortcuts

Undocumented commands:
=====
EOF eof exit q quit

Application commands (type help <topic>):
=====
=====
agent-delete          net-external-list    subnet-create
agent-list            net-gateway-connect  subnet-delete
agent-show            net-gateway-create   subnet-list
agent-update          net-gateway-delete   subnet-show
dhcp-agent-list-hosting-net net-gateway-disconnect subnet-update
dhcp-agent-network-add net-gateway-list
dhcp-agent-network-remove net-gateway-show
ext-list              net-gateway-update
ext-show              net-list
floatingip-associate  net-list-on-dhcp-agent
floatingip-create      net-show
floatingip-delete      net-update
floatingip-disassociate port-create
floatingip-list         port-delete
floatingip-show         port-list
help                   port-show
l3-agent-list-hosting-router port-update
l3-agent-router-add    queue-create
l3-agent-router-remove queue-delete
lb-healthmonitor-associate queue-list
lb-healthmonitor-create queue-show
lb-healthmonitor-delete quota-delete
lb-healthmonitor-disassociate quota-list
lb-healthmonitor-list   quota-show
lb-healthmonitor-show   quota-update
lb-healthmonitor-update router-create
lb-member-create        router-delete
lb-member-delete        router-gateway-clear
lb-member-list          router-gateway-set
lb-member-show          router-interface-add
lb-member-update        router-interface-delete
lb-pool-create          router-list
lb-pool-delete          router-list-on-l3-agent
lb-pool-list            router-port-list
lb-pool-show            router-show
lb-pool-stats           router-update
lb-pool-update          security-group-create
lb-vip-create           security-group-delete
lb-vip-delete           security-group-list
lb-vip-list             security-group-rule-create
lb-vip-show             security-group-rule-delete
lb-vip-update           security-group-rule-list
net-create              security-group-rule-show
net-delete              security-group-show

```

```
(quantum) net-list
+-----+-----+-----+
+-----+-----+-----+
| id                      | name          | subnets      |
+-----+-----+-----+
| 11fc08b7-c3b2-4b0c-bd04-66e279d9c470 | public_net1   | 13cc61f6-b33b-495a-a49f-83bdc9e439ab |
| 22f53ed1-3f3d-49c7-9162-7ba94d9c0a7e | private_mynet1 | b5a9b952-dd4f-445a-89c5-f15d0707b8bd |
| 2a405f54-aea0-47d7-8a43-4d5129e22b35 | test1         |                |
| d322e1ae-e068-4249-b9b3-7ed8b820bfa2 | mynetwork     |                |
+-----+-----+-----+
+-----+-----+-----+
```

Output format

We can use `-h` after each command to show the usage of each command:

```
(quantum) net-list -h
usage: quantum net-list [-h] [-f {csv,html,json,table,yaml}] [-c COLUMN]
                        [--quote {all,minimal,none,nonnumeric}]
                        [--request-format {json,xml}] [-D] [-F FIELD]
                        [-P SIZE] [--sort-key FIELD] [--sort-dir {asc,desc}]

List networks that belong to a given tenant.

optional arguments:
  -h, --help                show this help message and exit
  --request-format {json,xml}
                           the xml or json request format
  -D, --show-details        show detailed info
  -F FIELD, --field FIELD
                           specify the field(s) to be returned by server, can be
                           repeated
  -P SIZE, --page-size SIZE
                           specify retrieve unit of each request, then split one
                           request to several requests
  --sort-key FIELD           sort list by specified fields (This option can be
                           repeated), The number of sort_dir and sort_key should
                           match each other, more sort_dir specified will be
                           omitted, less will be filled with asc as default
                           direction
  --sort-dir {asc,desc}     sort list in specified directions (This option can be
                           repeated)

output formatters:
  output formatter options

  -f {csv,html,json,table,yaml}, --format {csv,html,json,table,yaml}
                           the output format, defaults to table
  -c COLUMN, --column COLUMN
                           specify the column(s) to include, can be repeated
```



```
CSV Formatter:
  --quote {all,minimal,none,nonnumeric}
           when to include quotes, defaults to nonnumeric
```

We can see the output formatters cliff provides to each command. By default, the output format is `table`. Now we choose `csv` output to run the command `net-list`:

```
(quantum) net-list -f csv
"id","name","subnets"
"11fc08b7-c3b2-4b0c-bd04-66e279d9c470","public_net1","13cc61f6-b33b-495a-a49f-83bdc9e439ab"
"22f53ed1-3f3d-49c7-9162-7ba94d9c0a7e","private_mynet1","b5a9b952-dd4f-445a-89c5-f15d0707b8bd"
"2a405f54-aea0-47d7-8a43-4d5129e22b35","test1",""
"d322e1ae-e068-4249-b9b3-7ed8b820bfa2","mynetwork",""
```

Column selection

We can see `-c COLUMN` in previous usage output. It can be used to limit the output fields:

```
(quantum) net-list -c id -c name
+-----+-----+
| id | name |
+-----+-----+
| 11fc08b7-c3b2-4b0c-bd04-66e279d9c470 | public_net1 |
| 22f53ed1-3f3d-49c7-9162-7ba94d9c0a7e | private_mynet1 |
| 2a405f54-aea0-47d7-8a43-4d5129e22b35 | test1 |
| d322e1ae-e068-4249-b9b3-7ed8b820bfa2 | mynetwork |
+-----+-----+
```

Features from API

Fields selection

If there are `'fields'` in request URL, V2.0 API will extract the list of fields to return.

A sample of such URLs is `http://localhost:9696/v2.0/networks.json?fields=id&fields=name`

Quantum client supports this feature by `-F` option in known options part and `--fields` in unknown options part. For example, `quantum -F id net-list -- --fields name`. Only `xx-list` and `xx-show` commands support this feature.

Value filtering

Any other fields except the `'fields'` are used to filter resources. A sample of such URLs is `http://localhost:9696/v2.0/networks.json?name=test1&name=test2`. By the current quantum server's sample DB plugin, the filtering has the same meaning as a SQL clause: `name in ['test1', 'test2']`. Quantum client supports this feature by any key options in unknown option part. For example `quantum net-list -- --name test1 test2`. Only `xx-list` commands support this feature.

Sample quantum command

All commands are run with following environment variables set:

```
export OS_USERNAME=admin
export OS_PASSWORD=password
export OS_TENANT_NAME=admin
export OS_AUTH_URL=http://localhost:5000/v2.0
```

- List the extensions of the system:

```
$ quantum ext-list -c alias -c name
+-----+-----+
| alias          | name          |
+-----+-----+
| agent_scheduler | Agent Schedulers |
| binding         | Port Binding    |
| quotas         | Quota management support |
| agent          | agent          |
| provider        | Provider Network |
| router          | Quantum L3 Router |
| lbaas           | LoadBalancing service |
| extraroute      | Quantum Extra Route |
+-----+-----+
```

- Create a network:

```
$ quantum net-create net1
Created a new network:
+-----+-----+
| Field          | Value          |
+-----+-----+
| admin_state_up | True           |
| id             | 2d627131-c841-4e3a-ace6-f2dd75773b6d |
| name           | net1           |
| provider:network_type | vlan           |
| provider:physical_network | physnet1       |
| provider:segmentation_id | 1001           |
| router:external | False          |
| shared         | False          |
| status         | ACTIVE         |
| subnets       |                |
| tenant_id      | 3671f46ec35e4bbca6ef92ab7975e463 |
+-----+-----+
```

Note: Some fields of the created network are invisible to non-admin users.

- Create a network with specified provider network type:

```
$ quantum net-create net2 --provider:network-type local
Created a new network:
+-----+-----+
| Field          | Value          |
+-----+-----+
```

admin_state_up	True
id	524e26ea-fad4-4bb0-b504-1ad0dc770e7a
name	net2
provider:network_type	local
provider:physical_network	
provider:segmentation_id	
router:external	False
shared	False
status	ACTIVE
subnets	
tenant_id	3671f46ec35e4bbca6ef92ab7975e463

Just as shown above, the unknown option `--provider:network-type` is used to create a local provider network.

- Create a subnet:

```
$ quantum subnet-create net1 192.168.2.0/24 --name subnet1
Created a new subnet:
```

Field	Value
allocation_pools	{"start": "192.168.2.2", "end": "192.168.2.254"}
cidr	192.168.2.0/24
dns_nameservers	
enable_dhcp	True
gateway_ip	192.168.2.1
host_routes	
id	15a09f6c-87a5-4d14-b2cf-03d97cd4b456
ip_version	4
name	subnet1
network_id	2d627131-c841-4e3a-ace6-f2dd75773b6d
tenant_id	3671f46ec35e4bbca6ef92ab7975e463

In the above command line, `net1` is the network name, `192.168.2.0/24` is the subnet's CIDR. They are positional arguments. `--name subnet1` is an unknown option, which specifies the subnet's name.

- Create a port with specified IP address:

```
$ quantum port-create net1 --fixed-ip ip_address=192.168.2.40
Created a new port:
```

Field	Value
admin_state_up	True
binding:capabilities	{"port_filter": false}

```

| binding:vif_type      | ovs
| device_id             |
| device_owner          |
| fixed_ips             | {"subnet_id": "15a09f6c-87a5-4d14-
b2cf-03d97cd4b456", "ip_address": "192.168.2.40"} |
| id                    | f7a08fe4-e79e-4b67-bbb8-a5002455a493
| mac_address           | fa:16:3e:97:e0:fc
| name                  |
| network_id            | 2d627131-c841-4e3a-ace6-f2dd75773b6d
| status                | DOWN
| tenant_id             | 3671f46ec35e4bbca6ef92ab7975e463
+-----+
+-----+
+

```

In the above command line, `net1` is the network name, which is a positional argument. `--fixed-ip ip_address=192.168.2.40` is an option, which specifies the port's fixed IP address we wanted.

- Create a port without specified IP address:

```

$ quantum port-create net1
Created a new port:
+-----+
+-----+
+
| Field                | Value
+-----+
+-----+
| admin_state_up       | True
| binding:capabilities | {"port_filter": false}
| binding:vif_type     | ovs
| device_id            |
| device_owner         |
| fixed_ips            | {"subnet_id": "15a09f6c-87a5-4d14-
b2cf-03d97cd4b456", "ip_address": "192.168.2.2"} |
| id                   | baf13412-2641-4183-9533-de8f5b91444c
| mac_address          | fa:16:3e:f6:ec:c7
| name                 |

```

```

| network_id      | 2d627131-c841-4e3a-ace6-f2dd75773b6d
| status         | DOWN
| tenant_id      | 3671f46ec35e4bbca6ef92ab7975e463
+-----+
+-----+
+

```

We can see that the system will allocate one IP address if we don't specify the IP address in command line.

- Query ports with specified fixed IP addresses:

```

$ quantum port-list --fixed-ips ip_address=192.168.2.2 ip_address=192.168.2.40
+-----+-----+-----+-----+
+-----+-----+-----+-----+
+
| id                  | name | mac_address      |
| fixed_ips          |      |                  |
|                    |      |                  |
+-----+-----+-----+-----+
+-----+-----+-----+-----+
+
| baf13412-2641-4183-9533-de8f5b91444c |      | fa:16:3e:f6:ec:c7 |
| {"subnet_id": "15a09f6c-87a5-4d14-b2cf-03d97cd4b456", "ip_address": "192.168.2.2"} |      |                  |
| f7a08fe4-e79e-4b67-bbb8-a5002455a493 |      | fa:16:3e:97:e0:fc |
| {"subnet_id": "15a09f6c-87a5-4d14-b2cf-03d97cd4b456", "ip_address": "192.168.2.40"} |      |                  |
+-----+-----+-----+-----+
+-----+-----+-----+-----+
+

```

--fixed-ips ip_address=192.168.2.2 ip_address=192.168.2.40 is one unknown option.

How to find unknown options? The unknown options can be easily found by watching the output of `create_xxx` or `show_xxx` command. For example, in the port creation command, we see the `fixed_ips` fields, which can be used as an unknown option.

quantum Client Reference

agent-delete	Delete a given agent.
agent-list	List agents.
agent-show	Show information of a given agent.
agent-update	Update a given agent.
dhcp-agent-list-hosting-net	List DHCP agents hosting a network.
dhcp-agent-network-add	Add a network to a DHCP agent.
dhcp-agent-network-remove	Remove a network from a DHCP agent.
ext-list	List all exts.
ext-show	Show information of a given resource

floatingip-associate fixed ip.	Create a mapping between a floating ip and a fixed ip.
floatingip-create	Create a floating ip for a given tenant.
floatingip-delete	Delete a given floating ip.
floatingip-disassociate fixed ip.	Remove a mapping from a floating ip to a fixed ip.
floatingip-list tenant.	List floating ips that belong to a given tenant.
floatingip-show	Show information of a given floating ip.
help	print detailed help for another command
l3-agent-list-hosting-router	List L3 agents hosting a router.
l3-agent-router-add	Add a router to a L3 agent.
l3-agent-router-remove	Remove a router from a L3 agent.
lb-healthmonitor-associate a pool.	Create a mapping between a health monitor and a pool.
lb-healthmonitor-create	Create a healthmonitor
lb-healthmonitor-delete	Delete a given healthmonitor.
lb-healthmonitor-disassociate pool.	Remove a mapping from a health monitor to a pool.
lb-healthmonitor-list tenant.	List healthmonitors that belong to a given tenant.
lb-healthmonitor-show	Show information of a given healthmonitor.
lb-healthmonitor-update	Update a given healthmonitor.
lb-member-create	Create a member
lb-member-delete	Delete a given member.
lb-member-list	List members that belong to a given tenant.
lb-member-show	Show information of a given member.
lb-member-update	Update a given member.
lb-pool-create	Create a pool
lb-pool-delete	Delete a given pool.
lb-pool-list	List pools that belong to a given tenant.
lb-pool-show	Show information of a given pool.
lb-pool-stats	Retrieve stats for a given pool.
lb-pool-update	Update a given pool.
lb-vip-create	Create a vip
lb-vip-delete	Delete a given vip.
lb-vip-list	List vips that belong to a given tenant.
lb-vip-show	Show information of a given vip.
lb-vip-update	Update a given vip.
net-create	Create a network for a given tenant.
net-delete	Delete a given network.
net-external-list tenant	List external networks that belong to a given tenant
net-gateway-connect router.	Add an internal network interface to a router.
net-gateway-create	Create a network gateway.
net-gateway-delete	Delete a given network gateway.
net-gateway-disconnect	Remove a network from a network gateway.
net-gateway-list	List network gateways for a given tenant.
net-gateway-show	Show information of a given network gateway.
net-gateway-update	Update the name for a network gateway.
net-list	List networks that belong to a given tenant.
net-list-on-dhcp-agent	List the networks on a DHCP agent.
net-show	Show information of a given network.
net-update	Update network's information.
port-create	Create a port for a given tenant.
port-delete	Delete a given port.
port-list	List ports that belong to a given tenant.
port-show	Show information of a given port.
port-update	Update port's information.

queue-create	Create a queue.
queue-delete	Delete a given queue.
queue-list	List queues that belong to a given tenant.
queue-show	Show information of a given queue.
quota-delete	Delete defined quotas of a given tenant.
quota-list	List defined quotas of all tenants.
quota-show	Show quotas of a given tenant
quota-update	Define tenant's quotas not to use defaults.
router-create	Create a router for a given tenant.
router-delete	Delete a given router.
router-gateway-clear	Remove an external network gateway from a router.
router-gateway-set	Set the external network gateway for a router.
router-interface-add	Add an internal network interface to a router.
router-interface-delete	Remove an internal network interface from a router.
router-list	List routers that belong to a given tenant.
router-list-on-l3-agent	List the routers on a L3 agent.
router-port-list	List ports that belong to a given tenant,
with specified router	
router-show	Show information of a given router.
router-update	Update router's information.
security-group-create	Create a security group.
security-group-delete	Delete a given security group.
security-group-list	List security groups that belong to a given tenant.
security-group-rule-create	Create a security group rule.
security-group-rule-delete	Delete a given security group rule.
security-group-rule-list	List security group rules that belong to a given tenant.
security-group-rule-show	Show information of a given security group rule.
security-group-show	Show information of a given security group.
subnet-create	Create a subnet for a given tenant.
subnet-delete	Delete a given subnet.
subnet-list	List networks that belong to a given tenant.
subnet-show	Show information of a given subnet.
subnet-update	Update subnet's information.

9. nova Client

Table of Contents

List Instances, Images, and Flavors	29
Launch an Instance	29
Commands Used	29
Before Launch	30
Create Your Server with the nova Client	30
Launch from a Volume	30
Associating ssh keys with instances	32
Insert metadata during launch	33
Providing User Data to Instances	34
Injecting Files into Instances	34
Change Server Configuration	34
Commands Used	34
Increase or Decrease Server Size	34
Instance evacuation	36
Before Evacuation	36
To evacuate your server without shared storage:	37
Evacuate server to specified host and preserve user data	37
Stop and Start an Instance	37
Pause and Unpause	37
Suspend and Resume	38
Rebooting an instance	38
Manage Security Groups	38
Add or delete a security group	38
Modify security group rules	39
Manage Floating IP Addresses	41
Reserve and associate floating IP addresses	41
Remove and de-allocate a floating IP address	42
Manage Images	43
Manage Volumes	43
Terminate an Instance	44
Get an Instance Console	44
Managing Bare metal Nodes	44
Command List for nova Client	46
Usage statistics	51
Host usage statistics	51
Instance usage statistics	52

This chapter describes how to use the nova client.

To install the client, see [Chapter 2, “Install the OpenStack Clients” \[3\]](#).

List Instances, Images, and Flavors

Before you can go about the business of building your cloud, you want to know what images are available to you by asking the image service what kinds of configurations are available. The image service could be compared to iTunes for your cloud - you can view the playlist of images before using your favorite image to create a new instance in the cloud. To get the list of images, their names, status, and ID, use this command:

```
$ nova image-list
```

ID	Name	Status	Server
53b205cc-7abc-46eb-aa60-eabc449b4217	natty-image	ACTIVE	
588d93af-645d-4312-a5b0-81347715a91b	tty-image	ACTIVE	
ac6f83b7-078c-47bd-b4c2-4053282da49e	oneiric-image	ACTIVE	
e110fb7d-2a9e-4da5-923f-5565867ce87a	maverick-image	ACTIVE	

Next you need to know the relative sizes of each of these.

```
$ nova flavor-list
```

ID	Name	Memory_MB	Disk	Ephemeral	Swap	VCPUs	RXTX_Factor
1	m1.tiny	512	0	0		1	1.0
2	m1.small	2048	10	20		1	1.0
3	m1.medium	4096	10	40		2	1.0
4	m1.large	8192	10	80		4	1.0
5	m1.xlarge	16384	10	160		8	1.0

You can also narrow down the list by using `grep` to find only the CentOS images with a command like this:

```
$ nova image-list | grep 'natty'
```

12	natty-server-cloudimg-amd64-kernel	ACTIVE	
13	natty-server-cloudimg-amd64	ACTIVE	

Launch an Instance

Launching a new instance on OpenStack.

Commands Used

This process uses the following commands:

- `nova boot`
- `nova list`
- `nova show`

Before Launch

With the information about what is available to you, you can choose the combination of image and flavor to create your virtual servers and launch instances.

Create Your Server with the nova Client

Procedure 9.1. To create and boot your server with the nova client:

1. Issue the following command. In the command, specify the server name, flavor ID, and image ID:

```
$ nova boot myUbuntuServer --image "3afe97b2-26dc-49c5-a2cc-a2fc8d80c001" --flavor 6
```

The command returns a list of server properties. The status field indicates whether the server is being built or is active. A status of `BUILD` indicates that your server is being built.

Property	Value
OS-DCF:diskConfig	AUTO
accessIPv4	
accessIPv6	
adminPass	ZbaYPZf6r2an
config_drive	
created	2012-07-27T19:59:31Z
flavor	8GB Standard Instance
hostId	
id	d8093de0-850f-4513-b202-7979de6c0d55
image	Ubuntu 12.04
metadata	{}
name	myUbuntuServer
progress	0
status	BUILD
tenant_id	345789
updated	2012-07-27T19:59:31Z
user_id	170454

2. Copy the server ID value from the `id` field in the output. You use this ID to get details for your server to determine if it built successfully.

Copy the administrative password value from the `adminPass` field. You use this value to log into your server.

Launch from a Volume

The Compute service has support for booting an instance from a volume.

Manually Creating a Bootable Volume

To manually create a bootable volume, mount the volume to an existing instance, and then build a volume-backed image. Here is an example based on [exercises/boot_from_volume.sh](#). This example assumes that you have a running instance with a 1GB

volume mounted at `/dev/vdc`. These commands will make the mounted volume bootable using a CirrOS image. As root:

```
# mkfs.ext3 -b 1024 /dev/vdc 1048576
# mkdir /tmp/stage
# mount /dev/vdc /tmp/stage

# cd /tmp
# wget https://launchpad.net/cirros/trunk/0.3.0/+download/cirros-0.3.0-x86_64-rootfs.img.gz
# gunzip cirros-0.3.0-x86_64-rootfs.img.gz
# mkdir /tmp/cirros
# mount /tmp/cirros-0.3.0-x86_64-rootfs.img /tmp/cirros

# cp -pr /tmp/cirros/* /tmp/stage
# umount /tmp/cirros
# sync
# umount /tmp/stage
```

Detach the volume once you are done.

Creating a Bootable Volume from an Image

Cinder has the ability to create a bootable volume from an image stored in Glance.

```
# cinder create --image-id <image_id> --display-name my-bootable-vol <size>
```

This feature is also mirrored in Nova:

```
# nova volume-create --image-id <image_id> --display-name my-bootable-vol
<size>
```



Note

As of Grizzly, the following block storage drivers are compatible: iSCSI-based, LVM, and Ceph.

Make sure you configure Cinder with the relevant Glance options:

Table 9.1. List of configuration flags for NFS

Flag Name	Type	Default	Description
<code>glance_host</code>	Optional	<code>\$my_ip</code>	(StrOpt) default glance hostname or ip
<code>glance_port</code>	Optional	9292	(IntOpt) default glance port
<code>glance_api_servers</code>	Optional	<code>\$glance_host:\$glance_port</code>	(ListOpt) A list of the glance api servers available to cinder: ([hostname ip]:port) (list value)
<code>glance_api_version</code>	Optional	1	(IntOpt) default version of the glance api to use
<code>glance_num_retries</code>	Optional	0	(IntOpt) Number retries when downloading an image from glance
<code>glance_api_insecure</code>	Optional	false	(BoolOpt) Allow to perform insecure SSL (https) requests to glance

Booting an instance from the volume

To boot a new instance from the volume, use the **nova boot** command with the `--block_device_mapping` flag. The output for **nova help boot** shows the following documentation about this flag:

```
--block_device_mapping <dev_name=mapping>
    Block device mapping in the format <dev_name>=
    <id>:<type>:<size(GB)>:<delete_on_terminate>.
```

The command arguments are:

<code>dev_name</code>	A device name where the volume will be attached in the system at <code>/dev/dev_name</code> . This value is typically <code>vda</code> .
<code>id</code>	The ID of the volume to boot from, as shown in the output of nova volume-list .
<code>type</code>	This is either <code>snap</code> , which means that the volume was created from a snapshot, or anything other than <code>snap</code> (a blank string is valid). In the example above, the volume was not created from a snapshot, so we will leave this field blank in our example below.
<code>size (GB)</code>	The size of the volume, in GB. It is safe to leave this blank and have the Compute service infer the size.
<code>delete_on_terminate</code>	A boolean to indicate whether the volume should be deleted when the instance is terminated. True can be specified as <code>True</code> or <code>1</code> . False can be specified as <code>False</code> or <code>0</code> .



Note

Because of bug [#1163566](#), you must specify an image when booting from a volume in Horizon, even though this image will not be used.

The following example will attempt boot from volume on the command line with ID=13, it will not delete on terminate. Replace the `--key_name` with a valid keypair name:

```
$ nova boot --flavor 2 --key_name mykey --block_device_mapping vda=13:::0
boot-from-vol-test
```

Associating ssh keys with instances

Creating New Keys

The command:

```
$ nova keypair-add mykey > mykey.pem
```

will create a key named `mykey` which you can associate with instances. Save the file `mykey.pem` to a secure location as it will allow root access to instances the `mykey` key is associated with.

Uploading Existing Keys

The command:

```
$ nova keypair-add --pub-key mykey.pub mykey
```

will upload the existing public key `mykey.pub` and associate it with the name `mykey`. You will need to have the matching private key to access instances associated with this key.

Adding Keys to Your Instance

To associate a key with an instance on boot add `--key_name mykey` to your command line for example:

```
$ nova boot --image ubuntu-cloudimage --flavor 1 --key_name mykey
```

Insert metadata during launch

When booting a server, you can also add metadata, so that you can more easily identify it amongst your ever-growing elastic cloud. Use the `--meta` option with a key=value pair, where you can make up the string for both the key and the value. For example, you could add a description and also the creator of the server.

```
$ nova boot --image=natty-image --flavor=2 smallimage2 --meta description='Small test image' --meta creator=joecool
```

When viewing the server information, you can see the metadata included on the metadata line:

```
$ nova show smallimage2
```

Property	Value
OS-DCF:diskConfig	MANUAL
OS-EXT-STS:power_state	1
OS-EXT-STS:task_state	None
OS-EXT-STS:vm_state	active
accessIPv4	
accessIPv6	
config_drive	
created	2012-05-16T20:48:23Z
flavor	m1.small
hostId	de0c201e62be88c61aeb52f51d91e147acf6cf2012bb57892e528487
id	8ec95524-7f43-4cce-a754-d3e5075bf915
image	natty-image
key_name	
metadata	{u'description': u'Small test image', u'creator': u'joecool'}
name	smallimage2
private network	172.16.101.11
progress	0
public network	10.4.113.11

status	ACTIVE
tenant_id	e830c2fbb7aa4586adf16d61c9b7e482
updated	2012-05-16T20:48:35Z
user_id	de3f4e99637743c7b6d27faca4b800a9
+-----+	

Providing User Data to Instances

User Data is a special key in the metadata service which holds a file that cloud aware applications within the guest instance can access. For example the [cloudinit](#) system is an open source package from Ubuntu that handles early initialization of a cloud instance that makes use of this user data.

This user-data can be put in a file on your local system and then passed in at instance creation with the flag `--user-data <user-data-file>` for example:

```
$ nova boot --image ubuntu-cloudimage --flavor 1 --user-data mydata.file
```

Injecting Files into Instances

Arbitrary local files can also be placed into the instance file system at creation time using the `--file <dst-path=src-path>` option. You may store up to 5 files. For example if you have a special `authorized_keys` file named `special_authorized_keysfile` that you want to put on the instance rather than using the regular [ssh key injection](#) for some reason you can use the following command:

```
$nova boot --image ubuntu-cloudimage --flavor 1 --file /root/.ssh/
authorized_keys=special_authorized_keysfile
```

Change Server Configuration

After you have created a server, you may need to increase its size, change the image used to build it, or perform other configuration changes.

Commands Used

This process uses the following commands:

- `nova resize*`
- `nova rebuild`

Increase or Decrease Server Size

Server size is changed by applying a different flavor to the server. Before you begin, use `nova flavor-list` to review the flavors available to you.

```
$ nova flavor-list
+-----+-----+-----+-----+-----+-----+-----+
| ID | Name | Memory_MB | Disk | Ephemeral | Swap | VCPUs | RXTX_Factor |
```

1	m1.tiny	512	0	0	1	1.0
2	m1.small	2048	10	20	1	1.0
3	m1.medium	4096	10	40	2	1.0
4	m1.large	8192	10	80	4	1.0
5	m1.xlarge	16384	10	160	8	1.0

In this example, we'll take a server originally configured with the `m1.tiny` flavor and resize it to `m1.small`.

```
$ nova show acdfb2c4-38e6-49a9-ae1c-50182fc47e35
```

Property	Value
OS-DCF:diskConfig	MANUAL
OS-EXT-STS:power_state	1
OS-EXT-STS:task_state	None
OS-EXT-STS:vm_state	active
accessIPv4	
accessIPv6	
config_drive	
created	2012-05-09T15:47:48Z
flavor	m1.tiny
hostId	
id	de0c201e62be88c61aeb52f51d91e147acf6cf2012bb57892e528487 acdfb2c4-38e6-49a9-ae1c-50182fc47e35
image	maverick-image
key_name	
metadata	{}
name	resize-demo
private network	172.16.101.6
progress	0
public network	10.4.113.6
status	ACTIVE
tenant_id	e830c2fbb7aa4586adf16d61c9b7e482

```

|         updated         |                2012-05-09T15:47:59Z
|         user_id         |                de3f4e99637743c7b6d27faca4b800a9
+-----+
+-----+

```

Use the `resize` command with the server's ID (6beefcf7-9de6-48b3-9ba9-e11b343189b3) and the ID of the desired flavor (2):

```
$ nova resize 6beefcf7-9de6-48b3-9ba9-e11b343189b3 2
```

While the server is rebuilding, its status will be displayed as `RESIZING`.

```

$ nova list
+-----+-----+-----+-----+
+-----+
| ID                               | Name           | Status  | Networks |
+-----+-----+-----+-----+
| 970e4ca0-f9b7-4c44-80ed-bf0152c96ae1 | resize-demo    | RESIZE  | private=172.16.101.6, public=10.4.113.6 |
+-----+-----+-----+-----+
+-----+

```

When the resize operation is completed, the status displayed is `VERIFY_RESIZE`. This prompts the user to verify that the operation has been successful; to confirm:

```
$ nova resize-confirm 6beefcf7-9de6-48b3-9ba9-e11b343189b3
```

However, if the operation has not worked as expected, you can revert it by doing:

```
$ nova resize-revert 6beefcf7-9de6-48b3-9ba9-e11b343189b3
```

In both cases, the server status should go back to `ACTIVE`.

Instance evacuation

As cloud administrator, while you are managing your cloud, you may get to the point where one of the cloud compute nodes fails. For example, due to hardware malfunction. At that point you may use server evacuation in order to make managed instances available again.

Before Evacuation

With the information about instance configuration, like if it is running on shared storage, you can choose the required evacuation parameters for your case. Use the **nova host-**

list command to list the hosts and find new host for the evacuated instance. In order to preserve user data on server disk, target host has to have preconfigured shared storage with down host. As well, you have to validate that the current vm host is down. Otherwise the evacuation will fail with error.

To evacuate your server without shared storage:

nova evacuate performs an instance evacuation from down host to specified host. The instance will be booted from a new disk, but will preserve the configuration, e.g. id, name, uid, ip...etc. New instance password can be passed to the command using the `--password <pwd>` option. If not given it will be generated and printed after the command finishes successfully.

```
$nova evacuate evacuated_server_name host_b
```

The command returns a new server password.

```
+-----+-----+
| Property | Value |
+-----+-----+
| adminPass | kRAJpErnT4xZ |
+-----+-----+
```

Evacuate server to specified host and preserve user data

In order to preserve the user disk data on the evacuated server the OpenStack Compute should be deployed with shared filesystem. Refer to the shared storage section in the [Configure migrations guide](#) in order to configure your system. In this scenario the password will remain unchanged.

```
$nova evacuate evacuated_server_name host_b --on-shared-storage
```

Stop and Start an Instance

There are two methods for stopping and starting an instance:

- **nova pause / nova unpause**
- **nova suspend / nova resume**

Pause and Unpause

nova pause stores the state of the VM in RAM. A paused instance continues to run, albeit in a "frozen" state.

Suspend and Resume

nova suspend initiates a hypervisor-level suspend operation. Suspending an instance stores the state of the VM on disk; all memory is written to disk and the virtual machine is stopped. Suspending an instance is thus similar to placing a device in hibernation, and makes memory and vCPUs available. Administrators may want to suspend an instance for system maintenance, or if the instance is not frequently used.

Rebooting an instance

nova reboot performs a reboot of a running instance. By default, this is a "soft" reboot, which will attempt a graceful shutdown and restart of the instance. To perform a "hard" reboot (i.e., a power cycle of the instance), pass the `--hard` flag as an argument.

Manage Security Groups

A security group is a named collection of network access rules that can be used to limit the types of traffic that have access to instances. When you spawn an instance, you can assign it to one or more groups. For each security group, the associated rules permit you to manage the allowed traffic to instances within the group. Any incoming traffic which is not matched by a rule is denied by default. At any time, it is possible to add or remove rules within a security group. Rules are automatically enforced as soon as they are created.

Before you begin, use **nova secgroup-list** to view the available security groups (specify `--all-tenants` if you are a cloud administrator wanting to view all tenants' groups) . You can also view the rules for a security group with **nova secgroup-list-rules**.

```
$ nova secgroup-list
+-----+-----+
| Name | Description |
+-----+-----+
| default | default |
+-----+-----+

$ nova secgroup-list-rules default
+-----+-----+-----+-----+-----+
| IP Protocol | From Port | To Port | IP Range | Source Group |
+-----+-----+-----+-----+-----+
| tcp | 80 | 80 | 0.0.0.0/0 | |
+-----+-----+-----+-----+-----+
```

In this example, the default security group has been modified to allow HTTP traffic on the instance by permitting TCP traffic on Port 80.

Add or delete a security group

Security groups can be added with **nova secgroup-create**.

The following example shows the creation of the security group `secure1`. After the group is created, it can be viewed in the security group list.

```
$ nova secgroup-create secure1 "Test security group"
+-----+
| Name | Description |
+-----+
| secure1 | Test security group |
+-----+

$ nova secgroup-list
+-----+
| Name | Description |
+-----+
| default | default |
| secure1 | Test security group |
+-----+
```

Security groups can be deleted with **nova secgroup-delete**. The default security group cannot be deleted. The default security group contains these initial settings:

- All the traffic originated by the instances (outbound traffic) is allowed
- All the traffic destined to instances (inbound traffic) is denied
- All the instances inside the group are allowed to talk to each other



Note

You can add extra rules into the default security group for handling the egress traffic. Rules are ingress only at this time.

In the following example, the group `secure1` is deleted. When you view the security group list, it no longer appears.

```
$ nova secgroup-delete secure1
$ nova secgroup-list
+-----+
| Name | Description |
+-----+
| default | default |
+-----+
```

Modify security group rules

The security group rules control the incoming traffic that is allowed to the instances in the group, while all outbound traffic is automatically allowed.



Note

It is not possible to change the default outbound behaviour.

Every security group rule is a policy which allows you to specify inbound connections that are allowed to access the instance, by source address, destination port and IP protocol, (TCP, UDP or ICMP). Currently, ipv6 and other protocols cannot be managed with the security rules, making them permitted by default. To manage such, you can deploy a firewall in front of your OpenStack cloud to control other types of traffic. The command requires the following arguments for both TCP and UDP rules :

- <secgroup> ID of security group.
- <ip_proto> IP protocol (icmp, tcp, udp).
- <from_port> Port at start of range.
- <to_port> Port at end of range.
- <cidr> CIDR for address range.

For ICMP rules, instead of specifying a begin and end port, you specify the allowed ICMP code and ICMP type:

- <secgroup> ID of security group.
- <ip_proto> IP protocol (with icmp specified).
- <ICMP_code> The ICMP code.
- <ICMP_type> The ICMP type.
- <cidr> CIDR for the source address range.



Note

Entering "-1" for both code and type indicates that all ICMP codes and types should be allowed.



The CIDR notation

That notation allows you to specify a base IP address and a suffix that designates the number of significant bits in the IP address used to identify the network. For example, by specifying a 88.170.60.32/27, you specify 88.170.60.32 as the **base IP** and 27 as the **suffix**. Since you use an IPV4 format, there are only 5 bits available for the host part (32 minus 27). The 0.0.0.0/0 notation means you allow the entire IPV4 range, meaning allowing all addresses.

For example, in order to allow any IP address to access to a web server running on one of your instance inside the default security group:

```
$ nova secgroup-add-rule default tcp 80 80 0.0.0.0/0
```

IP Protocol	From Port	To Port	IP Range	Source Group
tcp	80	80	0.0.0.0/0	

In order to allow any IP address to ping an instance inside the default security group (Code 0, Type 8 for the ECHO request.):

```
$ nova secgroup-add-rule default icmp 0 8 0.0.0.0/0
```

IP Protocol	From Port	To Port	IP Range	Source Group
icmp	0	8	0.0.0.0/0	

```
+-----+-----+-----+-----+
| icmp   | 0      | 8      | 0.0.0.0/0 |
+-----+-----+-----+-----+
```

```
$ nova secgroup-list-rules default
```

```
+-----+-----+-----+-----+-----+
| IP Protocol | From Port | To Port | IP Range | Source Group |
+-----+-----+-----+-----+-----+
| tcp        | 80        | 80      | 0.0.0.0/0 |               |
| icmp       | 0         | 8       | 0.0.0.0/0 |               |
+-----+-----+-----+-----+-----+
```

In order to delete a rule, you need to specify the exact same arguments you used to create it:

- <secgroup> ID of security group.
- <ip_proto> IP protocol (icmp, tcp, udp).
- <from_port> Port at start of range.
- <to_port> Port at end of range.
- <cidr> CIDR for address range.

```
$ nova secgroup-delete-rule default tcp 80 80 0.0.0.0/0
```

Manage Floating IP Addresses

A floating IP address is an IP address (typically public) that can be dynamically assigned to an instance. Pools of floating IP addresses are created outside of python-novaclient with the **nova-manage floating *** commands. Refer to "Configuring Public (Floating) IP Addresses" in the *OpenStack Compute Administration Manual* for more information.

Before you begin, use **nova floating-ip-pool-list** to determine what floating IP pools are available.

```
$ nova floating-ip-pool-list
+-----+
| name |
+-----+
| nova |
+-----+
```

In this example, the only available pool is nova.

Reserve and associate floating IP addresses

You can reserve floating IP addresses with the **nova floating-ip-create** command. This command reserves the addresses for the tenant, but does not immediately associate that address with an instance.

```
$ nova floating-ip-create nova
```

Ip	Instance Id	Fixed Ip	Pool
50.56.12.232	None	None	nova

The floating IP address has been reserved, and can now be associated with an instance with the **nova add-floating-ip** command. For this example, we'll associate this IP address with an image called `smallimage`.

```
$ nova add-floating-ip smallimage 50.56.12.232
```

After the command is complete, you can confirm that the IP address has been associated with the **nova floating-ip-list** and **nova-list** commands.

```
$ nova floating-ip-list
```

Ip	Instance Id	Fixed Ip	Pool
50.56.12.232	542235df-8ba4-4d08-90c9-b79f5a77c04f	10.4.113.9	nova

```
$ nova list
```

ID	Name	Status
4bb825ea-ea43-4771-a574-ca86ab429dcb	tinyimage2	ACTIVE
542235df-8ba4-4d08-90c9-b79f5a77c04f	smallimage	ACTIVE

The first table shows that the 50.56.12.232 is now associated with the `smallimage` instance ID, and the second table shows the IP address included under `smallimage`'s public IP addresses.

Remove and de-allocate a floating IP address

To remove a floating IP address from an instance, use the **nova remove-floating-ip** command.

```
$ nova remove-floating-ip smallimage 50.56.12.232
```

After the command is complete, you can confirm that the IP address has been associated with the **nova floating-ip-list** and **nova-list** commands.

```
$ nova floating-ip-list
+-----+-----+-----+-----+
|      Ip      | Instance Id | Fixed Ip | Pool |
+-----+-----+-----+-----+
| 50.56.12.232 |      None   |    None  | nova |
+-----+-----+-----+-----+

$ nova list
+-----+-----+-----+-----+
|          ID          |      Name      | Status |
+-----+-----+-----+-----+
| Networks            |                |        |
+-----+-----+-----+-----+
| 4bb825ea-ea43-4771-a574-ca86ab429dcb | tinyimage2 | ACTIVE | public=10.4.113.6; private=172.16.101.6 |
| 542235df-8ba4-4d08-90c9-b79f5a77c04f | smallimage | ACTIVE | public=10.4.113.9; private=172.16.101.9 |
+-----+-----+-----+-----+
```

You can now de-allocate the floating IP address, returning it to the pool so that it can be used by another tenant.

```
$ nova floating-ip-delete 50.56.12.232
```

In this example, 50.56.12.232 was the only IP address allocated to this tenant. Running **nova floating-ip-list** after the de-allocation is complete will return no results.

Manage Images

Adding images and setting the access to them can be managed in Glance, but you can create images by taking a snapshot of a running instance and view available images, set or delete image metadata, and delete an image, using the nova CLI.

Manage Volumes

Depending on the setup of your cloud provider, they may give you an endpoint to use to manage volumes, or there may be an extension under the covers. In either case, you can use the nova CLI to manage volumes.

```
volume-attach      Attach a volume to a server.
volume-create      Add a new volume.
volume-delete      Remove a volume.
volume-detach      Detach a volume from a server.
volume-list        List all the volumes.
volume-show        Show details about a volume.
volume-snapshot-create
Add a new snapshot.
volume-snapshot-delete
Remove a snapshot.
volume-snapshot-list
List all the snapshots.
volume-snapshot-show
Show details about a snapshot.
```

```

volume-type-create  Create a new volume type.
volume-type-delete  Delete a specific flavor
volume-type-list     Print a list of available 'volume types'.

```

Terminate an Instance

When you no longer need an instance, use the **nova delete** command to terminate it. You can use the instance name or the ID string. You will not receive a notification indicating that the instance has been deleted, but if you run the **nova list** command, the instance will no longer appear in the list.

In this example, we will delete the instance `tinyimage`, which is experiencing an error condition.

```

$ nova list
+-----+-----+-----+-----+
+-----+-----+-----+-----+
|          ID          |      Name      | Status |
+-----+-----+-----+-----+
| Networks |               |        |
+-----+-----+-----+-----+
+-----+-----+-----+-----+
| 30ed8924-fla5-49c1-8944-b881446a6a51 | tinyimage | ERROR | public=10.4.113.11; private=172.16.101.11 |
+-----+-----+-----+-----+
| 4bb825ea-ea43-4771-a574-ca86ab429dcb | tinyimage2 | ACTIVE | public=10.4.113.6; private=172.16.101.6 |
+-----+-----+-----+-----+
| 542235df-8ba4-4d08-90c9-b79f5a77c04f | smallimage | ACTIVE | public=10.4.113.9; private=172.16.101.9 |
+-----+-----+-----+-----+
+-----+-----+-----+-----+
$ nova delete tinyimage
$ nova list
+-----+-----+-----+-----+
+-----+-----+-----+-----+
|          ID          |      Name      | Status |
+-----+-----+-----+-----+
| Networks |               |        |
+-----+-----+-----+-----+
+-----+-----+-----+-----+
| 4bb825ea-ea43-4771-a574-ca86ab429dcb | tinyimage2 | ACTIVE | public=10.4.113.6; private=172.16.101.6 |
+-----+-----+-----+-----+
| 542235df-8ba4-4d08-90c9-b79f5a77c04f | smallimage | ACTIVE | public=10.4.113.9; private=172.16.101.9 |
+-----+-----+-----+-----+
+-----+-----+-----+-----+

```

Get an Instance Console

When you need to get a VNC console directly to a server, you can use the `nova get-vnc-console` command to connect.

Managing Bare metal Nodes

If you are using the Bare metal driver, you must inform each Nova Compute host of the physical hardware that it should manage. This is done via the baremetal admin extension.

Create a node with **nova baremetal-node-create**, and then add network interface information to it with **nova baremetal-interface-add**. Nodes and interfaces can be listed and deleted. When a node is in use, its status includes the UUID of the Nova instance that is running on it.

```
$ nova baremetal-node-create --pm_address=1.2.3.4 --pm_user=ipmi --
pm_password=ipmi $(hostname -f) 1 512 10 aa:bb:cc:dd:ee:ff
```

Property	Value
instance_uuid	None
pm_address	1.2.3.4
interfaces	[]
prov_vlan_id	None
cpus	1
memory_mb	512
prov_mac_address	aa:bb:cc:dd:ee:ff
service_host	ubuntu
local_gb	10
id	1
pm_user	ipmi
terminal_port	None

```
$ nova baremetal-interface-add 1 aa:bb:cc:dd:ee:ff
```

Property	Value
datapath_id	0
id	1
port_no	0
address	aa:bb:cc:dd:ee:ff

```
$ nova baremetal-node-list
```

ID	Host	CPUs	Memory_MB	Disk_GB	MAC Address
VLAN	PM Address	PM Username	PM Password	Terminal Port	
1	ubuntu	1	512	10	aa:bb:cc:dd:ee:ff
None	1.2.3.4	ipmi			None

```
$ nova boot --image my-baremetal-image --flavor my-baremetal-flavor test
```

Property	Value
status	BUILD
id	cc302a8f-cd81-484b-89a8-b75eb3911b1b

... wait for instance to become active ...

```
$ nova baremetal-node-show 1
```

Property	Value
----------	-------

```

| instance_uuid      | cc302a8f-cd81-484b-89a8-b75eb3911b1b |
| pm_address         | 1.2.3.4                               |
| interfaces         |                                         |
|   [{u'datapath_id': u'0', u'id': 1, u'port_no': 0, u'address':
|   u'aa:bb:cc:dd:ee:ff'}]} |
| prov_vlan_id       | None                                   |
| cpus                | 1                                     |
| memory_mb          | 512                                   |
| prov_mac_address    | aa:bb:cc:dd:ee:ff                   |
| service_host        | ubuntu                               |
| local_gb            | 10                                    |
| id                  | 1                                     |
| pm_user             | ipmi                                  |
| terminal_port       | None                                  |
+-----+-----+

```

Command List for nova Client

Example 9.1. nova Usage

```

usage: nova [--version] [--debug] [--os-cache] [--timings]
            [--timeout <seconds>] [--os-username <auth-user-name>]
            [--os-password <auth-password>]
            [--os-tenant-name <auth-tenant-name>] [--os-auth-url <auth-url>]
            [--os-region-name <region-name>] [--os-auth-system <auth-system>]
            [--service-type <service-type>] [--service-name <service-name>]
            [--volume-service-name <volume-service-name>]
            [--endpoint-type <endpoint-type>]
            [--os-compute-api-version <compute-api-ver>]
            [--os-cacert <ca-certificate>] [--insecure]
            [--bypass-url <bypass-url>]
            <subcommand> ...

```

Command-line interface to the OpenStack Nova API.

Example 9.2. Positional arguments

```

<subcommand>
  absolute-limits      Print a list of absolute limits for a user
  actions              Retrieve server actions.
  add-fixed-ip          Add new IP address to network.
  add-floating-ip       Add a floating IP address to a server.
  add-secgroup          Add a Security Group to a server.
  aggregate-add-host    Add the host to the specified aggregate.
  aggregate-create      Create a new aggregate with the specified details.
  aggregate-delete      Delete the aggregate by its id.
  aggregate-details     Show details of the specified aggregate.
  aggregate-list        Print a list of all aggregates.
  aggregate-remove-host Remove the specified host from the specified
                        aggregate.
  aggregate-set-metadata Update the metadata associated with the aggregate.
  aggregate-update      Update the aggregate's name and optionally
                        availability zone.
  baremetal-interface-add Add a network interface to a baremetal node

```

```

baremetal-interface-list      List network interfaces associated with a baremetal
                               node
baremetal-interface-remove    Remove a network interface from a baremetal node
baremetal-node-create         Create a baremetal node
baremetal-node-delete         Remove a baremetal node and any associated interfaces
baremetal-node-list           Print a list of available baremetal nodes
baremetal-node-show           Show information about a baremetal node
boot                          Boot a new server.
cloudpipe-create              Create a cloudpipe instance for the given project
cloudpipe-list                Print a list of all cloudpipe instances.
cloudpipe-update              Update a cloudpipe instance
console-log                   Get console log output of a server.
credentials                   Show user credentials returned from auth
delete                         Immediately shut down and delete a server.
diagnostics                   Retrieve server diagnostics.
dns-create                    Create a DNS entry for domain, name and ip.
dns-create-private-domain     Create the specified DNS domain.
dns-create-public-domain      Create the specified DNS domain.
dns-delete                    Delete the specified DNS entry.
dns-delete-domain             Delete the specified DNS domain.
dns-domains                   Print a list of available dns domains.
dns-list                      List current DNS entries for domain and ip or domain
                               and name.
endpoints                     Discover endpoints that get returned from the
                               authenticate services
evacuate                      Evacuates server from failed host
fixed-ip-get                  Show information for a fixed IP
fixed-ip-reserve              Reserve a fixed IP
fixed-ip-unreserve            Unreserve fixed IP
flavor-create                 Create a new flavor
flavor-delete                 Delete a specific flavor
flavor-key                    Set or unset extra_spec for a flavor.
flavor-list                   Print a list of available 'flavors' (sizes of
                               servers).
flavor-show                   Show details about the given flavor.
floating-ip-create            Allocate a floating IP for the current tenant.
floating-ip-delete            De-allocate a floating IP.
floating-ip-list              List floating ips for this tenant.
floating-ip-pool-list         List all floating ip pools.

get-password                  Get password for a server.
get-spice-console              Get a spice console to a server.
get-vnc-console               Get a vnc console to a server.

```

```

host-action                   Perform a power action on a host.
host-describe                 Describe a specific host
host-list                     List all hosts by service
host-update                   Update host settings.
hypervisor-list               List hypervisors.
hypervisor-servers            List instances belonging to specific hypervisors.

```

hypervisor-show	Display the details of the specified hypervisor.
hypervisor-stats	Get hypervisor statistics over all compute nodes.
hypervisor-uptime	Display the uptime of the specified hypervisor.
image-create	Create a new image by taking a snapshot of a running server.
image-delete	Delete an image.
image-list	Print a list of available images to boot from.
image-meta	Set or Delete metadata on an image.
image-show	Show details about the given image.
keypair-add	Create a new key pair for use with instances
keypair-delete	Delete keypair by its id
keypair-list	Print a list of keypairs for a user
list	List active servers.
list-extensions	List available extensions
live-migration	Migrates a running instance to a new machine.
lock	Lock a server.
meta	Set or Delete metadata on a server.
migrate	Migrate a server.
network-list	Print a list of available networks.
network-show	Show details about the given network.
pause	Pause a server.
quota-class-show	List the quotas for a quota class.
quota-class-update	Update the quotas for a quota class.
quota-defaults	List the default quotas for a tenant.
quota-show	List the quotas for a tenant.
quota-update	Update the quotas for a tenant.
rate-limits	Print a list of rate limits for a user
reboot	Reboot a server.
rebuild	Shutdown, re-image, and re-boot a server.
remove-fixed-ip	Remove an IP address from a server.
remove-floating-ip	Remove a floating IP address from a server.
remove-secgroup	Remove a Security Group from a server.
rename	Rename a server.
rescue	Rescue a server.
reset-state	Reset the state of an instance
resize	Resize a server.
resize-confirm	Confirm a previous resize.
resize-revert	Revert a previous resize (and return to the previous VM).
resume	Resume a server.
root-password	Change the root password for a server.
secgroup-add-group-rule	Add a source group rule to a security group.
secgroup-add-rule	Add a rule to a security group.
secgroup-create	Create a security group.
secgroup-delete	Delete a security group.
secgroup-delete-group-rule	Delete a source group rule from a security group.
secgroup-delete-rule	Delete a rule from a security group.
secgroup-list	List security groups for the current tenant.
secgroup-list-rules	List rules for a security group.
service-disable	Enable the service
service-enable	Enable the service
service-list	Show a list of all running services. Filter by host
and	service name.

show	Show details about the given server.
ssh	SSH into a server.
start	Start a server.
stop	Stop a server.
suspend	Suspend a server.
unlock	Unlock a server.
unpause	Unpause a server.
unrescue	Unrescue a server.
usage-list	List usage data for all tenants
volume-attach	Attach a volume to a server.
volume-create	Add a new volume.
volume-delete	Remove a volume.
volume-detach	Detach a volume from a server.
volume-list	List all the volumes.
volume-show	Show details about a volume.
volume-snapshot-create	Add a new snapshot.
volume-snapshot-delete	Remove a snapshot.
volume-snapshot-list	List all the snapshots.
volume-snapshot-show	Show details about a snapshot.
volume-type-create	Create a new volume type.
volume-type-delete	Delete a specific flavor.
volume-type-list	Print a list of available 'volume types'.
x509-create-cert	Create x509 cert for a user in tenant.
x509-get-root-cert	Fetches the x509 root cert.
bash-completion	Prints all of the commands and options to stdout so that the
help	Display help about this program or one of its subcommands.
network	Show a network
network-create	Create a network
network-delete	Delete a network
network-list	List networks
network	Show a network
network-create	Create a network
network-delete	Delete a network
network-list	List networks
virtual-interface-create	Add a new virtual interface to an instance
virtual-interface-delete	Removes the specified virtual interface from an instance
virtual-interface-list	Add a new virtual interface to an instance
backup-schedule	Show or edit the backup schedule for a server.
backup-schedule-delete	Delete the backup schedule for a server.
backup	Backup a server.
service-config	List available service attributes.
service-details	List available service attributes.
service-disable	Disable a specified service.
service-enable	Enable a specified service.
service-list	List available service attributes.
service-servers	List available service attributes.
service-show	List available service attributes.
service-version	List available service attributes.

```

baremetal-interface-add
    Add a network interface to a baremetal node
baremetal-interface-list
    List network interfaces associated with a baremetal
    node
baremetal-interface-remove
    Remove a network interface from a baremetal node
baremetal-node-create
    Create a baremetal node
baremetal-node-delete
    Remove a baremetal node and any associated interfaces
baremetal-node-list
    Print a list of available baremetal nodes
baremetal-node-show
    Show information about a baremetal node
list-extensions
    List all the os-api extensions that are available.
net
    Show a network
net-create
    Create a network
net-delete
    Delete a network
net-list
    List networks

```

Example 9.3. Optional arguments

```

--version          show program's version number and exit
--debug            Print debugging output
--os-cache         Use the auth token cache.
--timings          Print call timing info
--timeout <seconds> Set HTTP call timeout (in seconds)
--os-username <auth-user-name>
                  Defaults to env[OS_USERNAME].
--os-password <auth-password>
                  Defaults to env[OS_PASSWORD].
--os-tenant-name <auth-tenant-name>
                  Defaults to env[OS_TENANT_NAME].
--os-auth-url <auth-url>
                  Defaults to env[OS_AUTH_URL].
--os-region-name <region-name>
                  Defaults to env[OS_REGION_NAME].
--os-auth-system <auth-system>
                  Defaults to env[OS_AUTH_SYSTEM].
--service-type <service-type>
                  Defaults to compute for most actions
--service-name <service-name>
                  Defaults to env[NOVA_SERVICE_NAME]
--volume-service-name <volume-service-name>
                  Defaults to env[NOVA_VOLUME_SERVICE_NAME]
--endpoint-type <endpoint-type>
                  Defaults to env[NOVA_ENDPOINT_TYPE] or publicURL.
--os-compute-api-version <compute-api-ver>
                  Accepts 1.1, defaults to
env[OS_COMPUTE_API_VERSION].
--os-cacert <ca-certificate>
                  Specify a CA bundle file to use in verifying a
TLS
  (https) server certificate. Defaults to env[OS_CACERT]
--insecure
  Explicitly allow novaclient to perform "insecure"
SSL
  (https) requests. The server's certificate will
not be

```

```

This                                verified against any certificate authorities.
--bypass-url <bypass-url>          option should be used with caution.
Catalog                             Use this API endpoint instead of the Service

```

Usage statistics

The nova command-line tool can provide some basic statistics on resource usage for hosts and instances.

For more sophisticated monitoring, see the [Ceilometer](#) project, which is currently under development. You may also wish to consider installing tools such as [Ganglia](#) or [Graphite](#) if you require access to more detailed data.

Host usage statistics

Use the **nova host-list** command to list the hosts and the nova-related services that are running on them:

```

$ nova host-list
+-----+-----+
| host_name | service |
+-----+-----+
| c2-compute-01 | compute |
| c2-compute-01 | network |
| c2-compute-02 | compute |
| c2-compute-02 | network |
| c2-compute-03 | compute |
| c2-compute-03 | network |
| c2-compute-04 | compute |
| c2-compute-04 | network |
| c2-controller-01 | cert |
| c2-controller-01 | consoleauth |
| c2-controller-01 | scheduler |
+-----+-----+

```

Use the **nova host-describe** command to retrieve a summary of resource usage of all of the instances running on the host. The "cpu" column is the sum of the virtual CPUs of all of the instances running on the host, the "memory_mb" column is the sum of the memory (in MB) allocated to the instances running on the hosts, and the "disk_gb" column is the sum of the root and ephemeral disk sizes (in GB) of the instances running on the hosts.

Note that these values are computed using only information about the flavors of the instances running on the hosts. This command does not query the CPU usage, memory usage, or hard disk usage of the physical host.

```

$ nova host-describe c2-compute-01
+-----+-----+-----+-----+-----+
| HOST | PROJECT | cpu | memory_mb | disk_gb |
+-----+-----+-----+-----+-----+
| c2-compute-01 | (total) | 24 | 96677 | 492 |
+-----+-----+-----+-----+-----+

```

c2-compute-01 (used_max) 2 2560 0
c2-compute-01 (used_now) 4 7168 0
c2-compute-01 f34d8f7170034280a42f6318d1a4af34 2 2560 0
+-----+-----+-----+-----+
+-----+

Instance usage statistics

Use the **nova diagnostics** command to retrieve CPU, memory, I/O and network statistics from an instance:

```
$ nova diagnostics ubuntu
```

Property	Value
cpu0_time	1138410000000
memory	524288
memory-actual	524288
memory-rss	591664
vda_errors	-1
vda_read	334864384
vda_read_req	13851
vda_write	2985382912
vda_write_req	177180
vnet4_rx	45381339
vnet4_rx_drop	0
vnet4_rx_errors	0
vnet4_rx_packets	106426
vnet4_tx	37513574
vnet4_tx_drop	0
vnet4_tx_errors	0
vnet4_tx_packets	162200

Use the **nova usage-list** command to get summary statistics for each tenant:

```
$ nova usage-list
```

Usage from 2012-10-10 to 2012-11-08:

Tenant ID Disk GB-Hours	Instances	RAM MB-Hours	CPU Hours
0eec5c34a7a24a7a8ddad27cb81d2706	8	240031.10	468.81
92a5d9c313424537b78ae3e42858fd4e	5	483568.64	236.12
f34d8f7170034280a42f6318d1a4af34	106	16888511.58	9182.88

10. cinder Client

Table of Contents

cinder Command Reference	53
--------------------------------	----

This chapter describes how to use the cinder client.

To install the client, see [Chapter 2, “Install the OpenStack Clients”](#) [3].

cinder Command Reference

Command-line interface to the OpenStack Volume API.

Positional arguments:

<subcommand>	
absolute-limits	Print a list of absolute limits for a user
create	Add a new volume.
credentials	Show user credentials returned from auth
delete	Remove a volume.
endpoints	Discover endpoints that get returned from the authenticate services
extra-specs-list	Print a list of current 'volume types and extra specs' (Admin Only).
list	List all the volumes.
quota-class-show	List the quotas for a quota class.
quota-class-update	Update the quotas for a quota class.
quota-defaults	List the default quotas for a tenant.
quota-show	List the quotas for a tenant.
quota-update	Update the quotas for a tenant.
rate-limits	Print a list of rate limits for a user
rename	Rename a volume.
show	Show details about a volume.
snapshot-create	Add a new snapshot.
snapshot-delete	Remove a snapshot.
snapshot-list	List all the snapshots.
snapshot-rename	Rename a snapshot.
snapshot-show	Show details about a snapshot.
type-create	Create a new volume type.
type-delete	Delete a specific volume type
type-key	Set or unset extra_spec for a volume type.
type-list	Print a list of available 'volume types'.
bash-completion	Prints all of the commands and options to stdout so that the
help	Display help about this program or one of its subcommands.
list-extensions	List all the os-api extensions that are available.

Optional arguments:

--debug	Print debugging output
--os-username <auth-user-name>	Defaults to env[OS_USERNAME].
--os-password <auth-password>	

```
Defaults to env[OS_PASSWORD].
--os-tenant-name <auth-tenant-name>
Defaults to env[OS_TENANT_NAME].
--os-auth-url <auth-url>
Defaults to env[OS_AUTH_URL].
--os-region-name <region-name>
Defaults to env[OS_REGION_NAME].
--service-type <service-type>
Defaults to compute for most actions
--service-name <service-name>
Defaults to env[CINDER_SERVICE_NAME]
--volume-service-name <volume-service-name>
Defaults to env[CINDER_VOLUME_SERVICE_NAME]
--endpoint-type <endpoint-type>
Defaults to env[CINDER_ENDPOINT_TYPE] or publicURL.
--os-volume-api-version <compute-api-ver>
Accepts 1, defaults to env[OS_VOLUME_API_VERSION].
--os-cacert <ca-certificate>
Specify a CA bundle file to use in verifying a TLS
(https) server certificate. Defaults to env[OS_CACERT]
--retries <retries> Number of retries.
```

See "cinder help COMMAND" for help on a specific command.

11. swift Client

Table of Contents

swift Command Reference	55
-------------------------------	----

This chapter describes how to use the swift client.

To install the client, see [Chapter 2, “Install the OpenStack Clients”](#) [3].

swift Command Reference

```
Usage: swift <command> [options] [args]

Commands:
  stat [container] [object]
    Displays information for the account, container, or object depending on
    the
    args given (if any).
  list [options] [container]
    Lists the containers for the account or the objects for a container. -p or
    --prefix is an option that will only list items beginning with that
    prefix.
    -d or --delimiter is option (for container listings only) that will roll
    up
    items with the given delimiter (see Cloud Files general documentation for
    what this means).
  upload [options] container file_or_directory [file_or_directory] [...]
    Uploads to the given container the files and directories specified by the
    remaining args. -c or --changed is an option that will only upload files
    that have changed since the last upload. -S <size> or --segment-size
    <size>
    and --leave-segments are options as well (see --help for more).
  post [options] [container] [object]
    Updates meta information for the account, container, or object depending
    on
    the args given. If the container is not found, it will be created
    automatically; but this is not true for accounts and objects. Containers
    also allow the -r (or --read-acl) and -w (or --write-acl) options. The -m
    or --meta option is allowed on all and used to define the user meta data
    items to set in the form Name:Value. This option can be repeated. Example:
    post -m Color:Blue -m Size:Large
  download --all OR download container [options] [object] [object] ...
    Downloads everything in the account (with --all), or everything in a
    container, or a list of objects depending on the args given. For a single
    object download, you may use the -o [--output] <filename> option to
    redirect the output to a specific file or if "-" then just redirect to
    stdout.
  delete [options] --all OR delete container [options] [object] [object] ...
    Deletes everything in the account (with --all), or everything in a
    container, or a list of objects depending on the args given. Segments of
    manifest objects will be deleted as well, unless you specify the
    --leave-segments option.
```

Example:

```
swift -A https://auth.api.rackspacecloud.com/v1.0 -U user -K key stat
```

12. heat Client

Table of Contents

Create Stack	57
List Stacks	57
Stack Details	58
Update Stack	59
Heat command reference	59

This chapter describes how to use the heat client. This service orchestrates multiple composite cloud applications using a REST API that emulates the AWS CloudFormation API.

To install the client, see [Chapter 2, “Install the OpenStack Clients”](#) [3].

Create Stack

The following command will create a stack, or template, from an [example template file](#):

```
$ heat stack-create mystack --template-file=/path/to/heat/templates/
WordPress_Single_Instance.template
--parameters="InstanceType=m1.large;DBUsername=wp;DBPassword=
verybadpassword;KeyName=heat_key;LinuxDistribution=F17"
```

id	stack_name	stack_status
4c712026-dcd5-4664-90b8-0915494c1332	mystack	CREATE_IN_PROGRESS

The values specified in `--parameters` will depend on what parameters are defined in the template. If the template file is hosted on a website then the URL can be specified with `--template-url` instead of `--template-file`

As a template file is being developed, it can be validated without having to attempt stack creation:

```
$ heat stack-create mystack --template-file=/path/to/heat/templates/
WordPress_Single_Instance.template
```

The response will contain an error message if validation failed.

List Stacks

To see what stacks are visible to the current user:

```
$ heat stack-list
```

```

+-----+-----+-----+-----+
+-----+
| id                               | stack_name   | stack_status   |
| creation_time                   |              |                |
+-----+-----+-----+-----+
| 4c712026-dcd5-4664-90b8-0915494c1332 | mystack      | CREATE_COMPLETE |
| 2013-04-03T23:22:08Z |              |                |
| 7edc7480-bda5-4e1c-9d5d-f567d3b6a050 | my-otherstack | CREATE_FAILED   |
| 2013-04-03T23:28:20Z |              |                |
+-----+-----+-----+-----+
+-----+

```

Stack Details

There are a number of commands to explore the state and history of a particular stack. First, to show the details of a stack:

```
$ heat stack-show mystack
```

A stack consists of a collection of resources. This command will list all the resources in a stack, including their current status:

```

$ heat resource-list mystack
+-----+-----+-----+-----+
+-----+
| logical_resource_id | resource_type   | resource_status | updated_time
|                    |                |                |
+-----+-----+-----+-----+
+-----+
| WikiDatabase        | AWS::EC2::Instance | CREATE_COMPLETE |
| 2013-04-03T23:25:56Z |                    |                |
+-----+-----+-----+-----+
+-----+

```

This command will show the details for the specified resource in a stack:

```
$ heat resource-show mystack WikiDatabase
```

Some resources have associated metadata which can change throughout the life-cycle of a resource:

```
$ heat resource-metadata mystack WikiDatabase
```

A series of events is generated during the life-cycle of a stack. This command will display those events.

```

$ heat event-list mystack
+-----+-----+-----+-----+
+-----+
| logical_resource_id | id | resource_status_reason | resource_status |
| event_time         |    |                        |                |
+-----+-----+-----+-----+
+-----+
| WikiDatabase        | 1  | state changed          | IN_PROGRESS     |
| 2013-04-03T23:22:09Z |    |                        |                |
+-----+-----+-----+-----+

```

```
| WikiDatabase | 2 | state changed | CREATE_COMPLETE |
2013-04-03T23:25:56Z |
+-----+-----+-----+-----+
+-----+
```

This command will show the details for a particular event:

```
$ heat event-show WikiDatabase 1
```

Update Stack

The following command is an example of updating an existing stack from a modified template file:

```
$ heat stack-update mystack --template-file=/path/to/heat/templates/
WordPress_Single_Instance_v2.template
--parameters="InstanceType=m1.large;DBUsername=wp;DBPassword=
verybadpassword;KeyName=heat_key;LinuxDistribution=F17"
+-----+-----+-----+-----+
+-----+
| id | stack_name | stack_status |
| creation_time |
+-----+-----+-----+-----+
| 4c712026-dcd5-4664-90b8-0915494c1332 | mystack | UPDATE_COMPLETE |
2013-04-03T23:22:08Z |
| 7edc7480-bda5-4e1c-9d5d-f567d3b6a050 | my-otherstack | CREATE_FAILED |
2013-04-03T23:28:20Z |
+-----+-----+-----+-----+
+-----+
```

Depending on what is being updated, some resources will be updated in-place, while others will be replaced with new resources.

Heat command reference

```
usage: heat [-d] [-v] [-k] [--cert-file CERT_FILE] [--key-file KEY_FILE]
           [--ca-file CA_FILE] [--timeout TIMEOUT]
           [--os-username OS_USERNAME] [--os-password OS_PASSWORD]
           [--os-tenant-id OS_TENANT_ID] [--os-tenant-name OS_TENANT_NAME]
           [--os-auth-url OS_AUTH_URL] [--os-region-name OS_REGION_NAME]
           [--os-auth-token OS_AUTH_TOKEN] [--os-no-client-auth]
           [--heat-url HEAT_URL] [--heat-api-version HEAT_API_VERSION]
           [--os-service-type OS_SERVICE_TYPE]
           [--os-endpoint-type OS_ENDPOINT_TYPE] [-t]
           <subcommand> ...

stack-create      Create the stack
stack-delete      Delete the stack
stack-list        List the user's stacks
stack-show        Describe the stack
stack-update      Update the stack
resource-list     Show list of resources belonging to a stack
resource-metadata List resource metadata
resource-show     Describe the resource
event-list        List events for a stack
event-show        Describe the event
```

```
template-show      Get the template for the specified stack
template-validate  Validate a template with parameters
help              Display help about this program or one of its
subcommands.
```

Optional arguments:

```
-d, --debug          Defaults to env[HEATCLIENT_DEBUG]
-v, --verbose        Print more verbose output
-k, --insecure        Explicitly allow the client to perform "insecure" SSL
                      (https) requests. The server's certificate will not be
                      verified against any certificate authorities. This
                      option should be used with caution.

--cert-file CERT_FILE Path of certificate file to use in SSL connection.
                      This file can optionally be prepended with the private
                      key.

--key-file KEY_FILE   Path of client key to use in SSL connection. This
                      option is not necessary if your key is prepended to
                      your cert file.

--ca-file CA_FILE     Path of CA SSL certificate(s) used to verify the
                      remote server's certificate. Without this option the
                      client looks for the default system CA certificates.

--timeout TIMEOUT     Number of seconds to wait for a response

--os-username OS_USERNAME
                      Defaults to env[OS_USERNAME]
--os-password OS_PASSWORD
                      Defaults to env[OS_PASSWORD]
--os-tenant-id OS_TENANT_ID
                      Defaults to env[OS_TENANT_ID]
--os-tenant-name OS_TENANT_NAME
                      Defaults to env[OS_TENANT_NAME]
--os-auth-url OS_AUTH_URL
                      Defaults to env[OS_AUTH_URL]
--os-region-name OS_REGION_NAME
                      Defaults to env[OS_REGION_NAME]
--os-auth-token OS_AUTH_TOKEN
                      Defaults to env[OS_AUTH_TOKEN]
--os-no-client-auth    Do not contact keystone for a token. Defaults to
                      env[OS_NO_CLIENT_AUTH]
--heat-url HEAT_URL    Defaults to env[HEAT_URL]
--heat-api-version HEAT_API_VERSION
                      Defaults to env[HEAT_API_VERSION] or 1
--os-service-type OS_SERVICE_TYPE
                      Defaults to env[OS_SERVICE_TYPE]
--os-endpoint-type OS_ENDPOINT_TYPE
                      Defaults to env[OS_ENDPOINT_TYPE]
-t, --token-only       Only send a token for auth, do not send username and
                      password as well.
```

See "heat help COMMAND" for help on a specific command.